


Decision Team

A harness for making consequential decisions well – build vs. buy, pricing, market entry, hiring, vendor selection. Six analytical roles run in isolation, every claim traces to evidence, an independent red team tries to kill the recommendation, and the deliverable is a list of cheap, falsifiable tests – not a persuasive memo.

HOW A DECISION MOVES THROUGH THE HARNESS



What a skill is – and how one is built

A skill is a folder of instructions, references, templates, and scripts that Claude loads on demand. Instead of re-explaining a process every session, the skill captures it once – the workflow, the guardrails, the output format – and Claude follows it consistently.

Anatomy of a skill

decision-team/

```

├─ SKILL.md ← required · frontmatter + playbook
├─ references/ ← depth, loaded per step
│   └─ roles.md
│   └─ evidence-protocol.md
│   └─ redteam-protocol.md
│   └─ memo-format.md
├─ assets/ ← templates + worked examples
│   └─ workspace/
│   └─ examples/
└─ scripts/ ← deterministic checks
    └─ init_workspace.py
    └─ evidence_lint.py
    └─ rank_assumptions.py

```

Progressive disclosure. Claude reads the playbook up front and pulls reference files only at the step that needs them. A skill that front-loads everything burns context on material that may never be used.

The three design rules that matter most

FRONTMATTER

The description is a trigger, not a summary

It's the only thing Claude sees before deciding to load the skill. Name the concrete situations (*build or buy, raise prices, kill a product*) and the casual phrasings ("should we...", "I'm torn between...") that should activate it.

PLAYBOOK

Explain why, put gates early, name failure modes

A model that understands the rationale applies rules correctly in situations the author didn't anticipate. If there's a condition under which the run should stop, check it first – and make stopping a success outcome.

SCRIPTS

Script the checks a model will fudge


Counting, thresholds, format validation, computed priority. Standard library only, runnable from a fresh clone – and the script should print what it *cannot* check, so those conditions get applied manually.

Why six voices aren't six perspectives

The obvious way to do multi-role analysis is to ask one model to play six characters in one conversation. That fails in a specific, invisible way – and the skill's entire structure exists to prevent it.

One conversation, six characters

✗ CORRELATED ERRORS



The errors are correlated, and synthesis cannot detect correlated error because it looks exactly like agreement. Consensus that *feels* like converging evidence is precisely the failure a decision process exists to prevent.

Isolated agents, sliced evidence

✓ INDEPENDENT ERRORS



Each role sees only its slice, its brief, and the decision statement – never the other roles' outputs. Where roles disagree, the disagreement arose from *different evidence*, which makes it information rather than theatre.

1

Isolation

Independence has to be manufactured structurally – it will not arise from prompting. Disjoint evidence slices, no sight of each other's reasoning.

2

Traceability

Every claim carries a tag: the evidence item that supports it, or an explicit admission that nothing does. The lint script makes fluent unsupported prose expensive.

3

Refutation, not review

The red team runs afterward, in a clean context, sees the conclusion and raw evidence but never the reasoning – and is asked to break the decision, not assess it.

Ten steps, four phases

Run in order. The gate can end the run at step 2 – and when it does, that is the skill working, not failing.

PHASE A · FRAME & GATE

STEP 1

Build workspace, state the decision

Interview the user – never infer. The move, the alternative, the date, observable success. Gaps marked [NOT SUPPLIED] are the first finding.

STEP 2

Load evidence, run the gate

< 5 primary / internal-data items – or demand-hinged with zero buyer conversations – **stop**. Deliver an Evidence Acquisition Plan instead.

STEP 3

Fill context, slice evidence

Map each evidence item to the roles that should see it. Each slice is a standalone file – never "the whole file, ignore parts."

PHASE B · ISOLATED ANALYSIS

STEP 4

Run five roles in parallel isolation

Market Analyst · Customer Strategist · Operator · Finance Lead · Historian. Each ends with "Evidence I was not given but need."

STEP 5

Lint each deliverable against its slice

The one automated check that isolation held: a role citing evidence it was never handed gets caught. Failures go back to the role – never patched by hand.

PHASE C · SYNTHESIS & ATTACK

STEP 6

Map disagreement mechanically

Extract claims on the same assumption; find polarity conflicts. Total agreement is a warning, not a green light.

STEP 7

Build the assumption ledger

Load-bearing? Importance × uncertainty. Flag every load-bearing assumption with no cheap test attached.

STEP 8

Draft memo...

Tagged synthesis linted at the 0.20 ceiling; the memo rendered from it in readable prose.

STEP 8b

...then red-team it

Three refuters, distinct lenses, clean context. A refutation only counts if it names a ~30-day observable.

PHASE D · VERDICT & TESTS

STEP 9

Set the verdict – constrained, not chosen

Proceed / modify / test / stop, bound to the ledger. No exception for a strong narrative.

STEP 10

Write the test cards

Metric, cost, sample – and pass / fail / modify thresholds committed **before** the test runs, timestamped.

The most-skipped, most-valuable part: the evidence gate. Running six roles on thin inputs produces a beautifully formatted hallucination with a financial model attached – more dangerous than no analysis, because it feels like diligence. A user who came for a strategy memo and left with three phone calls to make got the better deal.

The gate decides if analysis runs.

The ledger decides what it may conclude.

The evidence gate

Evidence loaded into 03-evidence.md
one block per item · E<n> id · source · type

≥ 5 items that are primary or internal-data,
from independent sources?

YES ↓

NO ↓

Decision hinges on demand
with zero buyer conversations?

YES ↓

NO ↓

Stop

Gate passes
proceed to slicing
& roles

Stop. No memo.

Deliver an **Evidence Acquisition Plan**:
the 3 cheapest things to find out
this week,
what each settles, and what
you'd do
differently depending on the
result.

Verdict constraints

The verdict is one of **proceed / modify / test / stop** – and it is constrained rather than chosen freely, because the pull toward an agreeable, confident "proceed" is strong and does not announce itself.

CONDITION IN THE LEDGER

FORCED VERDICT

Load-bearing assumption, priority
≥ 16, no completed test

TEST

≥ 2 surviving red-team refutations
on one load-bearing assumption

TEST /
STOP

Evidence gate passed only
marginally

TEST – and say
so

No forcing condition triggered

PROCEED is
allowed

Primary = you or the user talked to a real buyer, ran the test, shipped the thing. **Internal-data** = your own numbers, exported – not remembered. Web research is secondary and carries a URL. Nothing is ever written from the model's priors.

Either/or decisions get two verdicts. "A: test. B: proceed." is a real answer; collapsing it into a single token is not. If the framing itself is broken – both options fail, or they aren't actually exclusive – say that first.

What gets delivered, in order of prominence

1

The test list

Three to five cards, cheapest first, thresholds pre-committed. The primary deliverable.

2

The decision memo

One page, evidence-quality footer intact – the reader's only way to calibrate what they're holding.

3

The assumption ledger

Sorted by importance × uncertainty; untested load-bearing items flagged.

4

Open disagreements & surviving refutations

Unresolved, as they are. Writing smoothly over a real disagreement is the most expensive thing this process can do.

Adapting the template to other decisions

Nothing about the harness is specific to build-vs-buy. Keep the invariant core, swap the variant surface, and the same structure handles any consequential, hard-to-reverse call made under uncertainty.

Keep – the invariant core

Evidence gate – no analysis on thin inputs, ever

Isolated roles – independence is structural, not prompted

Claim traceability – every claim tagged or admitted

Refutation after synthesis – with a ~30-day observable

Ledger-bound verdict – not chosen by narrative strength

Swap – the variant surface

Trigger description – the decision types that activate it

Role cast – the five perspectives that matter here

"Primary evidence" – this domain's version of "talked to a real buyer"

Gate stop-condition – built around what the decision hinges on

Red-team lenses – how this kind of decision actually fails

Pricing decision

Recast: Pricing Analyst · Customer Strategist · Finance · Operator · Historian. **Primary evidence:** willingness-to-pay conversations; own churn data by segment.

Gate: hinges on elasticity + zero price conversations → stop.

Key hire / org change

Recast: Talent Analyst · Operator · Finance · Team Strategist · Historian. **Primary evidence:** references you actually called; measured cost of the gap.

Gate: "we're drowning" with no measurement → stop, measure first.

Market entry

Recast: Market Analyst · Customer Strategist · Operator · Finance · Historian. **Primary evidence:** buyers in the target market – home-market enthusiasm doesn't count.

Gate: zero target-market conversations → automatic stop.

Vendor selection

Recast: Technical Evaluator · Operator · Finance · Risk Analyst · Historian. **Primary evidence:** a PoC on your own workload; reference calls at your scale.

Gate: no PoC and no reference calls → stop. A sales demo is secondary.

Kill / sunset decision

Recast: Customer Strategist · Finance · Operator · Market Analyst · Historian. **Primary evidence:** per-account usage, exported; conversations with the heaviest users.

Gate: nobody has talked to the top users → stop.

Your variant

Copy the folder, rewrite the trigger, recast the roles, redefine primary evidence, re-aim the red team, add domain verdict rules, update the worked example, repackage.

Full guide: docs/creating-and-adapting-skills.md

The one rule that transfers everywhere: every variant exists to answer the same question – what is the short list of cheap, falsifiable tests that would actually change what you do? If an adaptation starts optimising for a more persuasive memo instead of a better test list, it has lost the plot.